

# AGNIBH DASGUPTA

✉ [adg002@gmail.com](mailto:adg002@gmail.com)

☎ +1 (435) 754-9599

📍 San Francisco, CA, USA

🌐 Website: [cent664.github.io](https://cent664.github.io)

🐙 Github: [github.com/cent664](https://github.com/cent664)

in LinkedIn: [linkedin.com/in/cent664](https://linkedin.com/in/cent664)

🎓 Google Scholar:

<https://tinyurl.com/cent-google scholar>

## RESEARCH INTERESTS

Representation learning, Watermarking, Multimodality, Provenance

## EDUCATION

**University of Nebraska Omaha**, Omaha, Nebraska, USA

Ph.D., Computing & Information Science

2021 – 2026

Dissertation: *Invariant Feature Learning in AI Models: Exploring Representation Spaces for Robust Watermarking and Beyond* (Dr. Xin Zhong)

**Utah State University**, Logan, Utah, USA

M.S., Computer Science

2018 – 2020

Thesis: *Deep Q-Learning Applied to Stock Trading* (Dr. Nicholas Flann)

**West Bengal University of Technology**, Kolkata, India

B.Tech., Computer Science

2013 – 2017

## PROFESSIONAL EXPERIENCE

**University of Nebraska Omaha**, Graduate Researcher

2021 – 2026

- Developed representation-level approaches for robust watermarking and model attribution across vision and language models, including semantically grounded invariant feature learning, black-box LLM watermarking, and geometric analysis of invariant subspaces in pretrained LLMs.
- Collaborated with medical practitioners on retinal vascular leakage detection from ultra-wide-field fluorescein angiography scans; trained patch-based classifiers and presented findings at the American Academy of Ophthalmology (AAO), 2025.

## Independent Projects

2026

- Designed and deployed Siosa's Library, a production agentic RAG system for grounded Q&A; the pipeline integrates a LangGraph planner, multi-step wiki retrieval, live poewiki search with passage reranking, and LLM-as-judge evaluation, shipped with FastAPI.

## PUBLICATIONS

### Conference Publications

- [C1] A. A. Tanvir, **A. Dasgupta**, X. Zhong. **TIACam: Text-Anchored Invariant Feature Learning with Auto-Augmentation for Camera-Robust Zero-Watermarking**. *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026.
- [C2] H. Yoon, X. Zhong, **A. Dasgupta**, G. Nugent, G. Trainin. **Leveraging Artificial Intelligence (AI) to Enhance Computer Science Instruction**. *IEEE Frontiers in Education Conference (FIE)*, 2024.
- [C3] **A. Dasgupta**, X. Zhong. **Robust Image Watermarking Based on Cross-Attention and Invariant Domain Learning**. *International Conference on Computational Science and Computational Intelligence (CSCI)*, 2023.
- [C4] N. Khatri, **A. Dasgupta**, Y. Shen, X. Zhong, F. Y. Shih. **Perspective Transformation Layer**. *International Conference on Computational Science and Computational Intelligence (CSCI)*, 2022.

## Journal Publications

- [J1] A. Dasgupta, A. A. Tanvir, X. Zhong. [Watermarking Language Models through Language Models](#). *IEEE Transactions on Artificial Intelligence*, pp. 1–10, 2025. DOI: 10.1109/TAI.2025.3605117

## Book Chapters

- [B1] A. Dasgupta, X. Zhong. [Enhanced image watermarking through cross-attention and noise-invariant domain learning](#). *Book Chapter, Imaging Science: Computer Vision, Image and Signal Processing, Pattern Recognition*, IEEE, 2024.

## Submitted / Under Review

- [S1] A. Dasgupta, A. A. Tanvir, X. Zhong. [Invariant Features in Language Models: Geometric Characterization and Model Attribution](#).

## THESES AND OTHER RESEARCH

---

- [O1] [Invariant Feature Learning in AI Models: Exploring Representation Spaces for Robust Watermarking and Beyond](#) 2026
- Doctoral dissertation investigating invariant representation learning across vision and language domains, using watermarking as a quantitative probe of representational stability.
- [O2] [Retinal Vascular Leakage Detection](#) 2025
- Trained classifiers on patches of ultra-wide-field fluorescein angiography scans to detect areas of leakage, in collaboration with medical practitioners. Presented at the American Academy of Ophthalmology (AAO), 2025.
- [O3] [Sentiment Analysis on Electric Vehicle Discourse](#) 2022
- Analyzed public sentiment using lexical and Word2Vec-LSTM approaches on scraped Twitter data.
- [O4] [Deep Q-Learning Applied to Stock Trading](#) 2020
- Designed and evaluated a reinforcement learning trading system using auto-scraped S&P 500 index data as part of the M.S. thesis.

## TEACHING EXPERIENCE

---

Responsibilities across all courses included grading, holding office hours, and lab supervision where applicable. For Capstone classes, I provided guidance on project scope, feasibility, and implementation strategy for senior undergraduate projects.

**University of Nebraska Omaha, Graduate Teaching Assistant** 2023–2026

- **CSCI1620: Introduction to Computer Science II** (Undergraduate), Fall 2023 (class size: 59), Fall 2024 (class size: 32)  
*Course overview:* Covers object-oriented programming, data structures, and algorithm design in an introductory CS sequence. Designed quizzes with progressive difficulty structures in Fall 2024.
- **CSCI4970: Capstone Project** (Undergraduate), Spring 2024 (class size: 67), Fall 2025 (class size: 58), Spring 2026 (class size: 90)  
*Course overview:* Integrative senior project course in which students design and deliver a substantial service-learning oriented software or research project.

**University of Nebraska Omaha, Intern Supervisor** MAY – AUG 2021

- Interviewed and shortlisted high school intern candidates for AI research positions at the [DLM Lab](#).

- Supervised interns on active research projects in representation learning and computer vision, adapting technical content for students with varying levels of prior knowledge.

Utah State University, Graduate Teaching Assistant

2018–2020

- **CS-4700: Programming Languages** (Undergraduate), Fall 2018 (class size: 49), Spring 2019 (class size: 54), Fall 2019 (class size: 64), Spring 2020 (class size: 69)  
*Course overview:* Covers paradigms, syntax, semantics, and implementation of programming languages including functional and logic programming. Designed assignments with a progressive difficulty structure in Spring 2020.
- **ECE-1410: Computer Programming for ECE II** (Undergraduate), Fall 2020 (class size: 26)  
*Course overview:* Continuation of introductory programming for electrical and computer engineering students, covering data structures and algorithmic problem solving.
- **CS-5890/6890: Temporal Data Mining** (Cross-listed), Fall 2020 (class size: 13)  
*Course overview:* Cross-listed course covering techniques for mining patterns and knowledge from time-series and sequential data.

## GRANTS AND AWARDS

**Graduate Research and Creative Activity (GRACA) Fellowship, \$5,000** 2025

University of Nebraska Omaha, Office of Research and Creative Activity

- Supported summer research on training a neural network augementer for automatic image augmentations; presented at UNO Research & Creative Activity Fair

**Graduate Research and Creative Activity (GRACA) Fellowship, \$5,000** 2021

University of Nebraska Omaha, Office of Research and Creative Activity

- Supported summer research on invariant domain learning for image watermarking; presented at UNO Research & Creative Activity Fair

## SERVICE

### External Reviewer

- ACCV, 2022
- Scientific Reports (Nature), 2023
- Springer Nature, 2025

## SOFTWARE AND CODE

**Siosa's Library (2026)** — <https://github.com/cent664/Siosa>

A production wiki-grounded RAG application for Path of Exile 1 mechanics Q&A, deployed at [poesiosa.net](https://poesiosa.net). At ask time the system searches poewiki.net (multiple query variants and title probes), fetches relevant pages, reranks passages, and generates short cited answers through swappable Claude/GPT-4 API backends. A LangGraph planner can issue several targeted wiki lookups for multi-topic questions before synthesis. Evaluation uses a hand-labeled gold set (page-level retrieval precision/recall) and optional LLM-as-judge scores for faithfulness and relevance; production disables inline judges for latency. Built with FastAPI and React (trace UI for development); containerized deploy on Railway with GitHub Actions CI.

**LLMID (2026)** — <https://github.com/cent664/LLMID>

A geometric framework for identifying invariant subspaces in pretrained LLMs across 9 open-source model families, without additional training. The framework applies contrastive sensitivity decomposition to separate semantic and nuisance directions in hidden representations, validated through causal patching

experiments. The discovered invariant zones achieve 90%+ zero-shot model attribution accuracy across base, fine-tuned, and distilled model variants.

**TIACam (2026)** — <https://github.com/tanvir097/TIACam>

A camera-robust zero-watermarking framework that combines CLIP image-text features with a learned adversarial augmentation module to ground invariant representations in semantic content. The auto-augmentor discovers camera-like distortions through differentiable operators, removing the need for manually designed augmentation families. Achieves 95%+ watermark recovery on ImageNet and Visual Genome datasets under real-world camera distortions.

**LLMWM (2025)** — <https://github.com/cent664/LLMWM>

A model-agnostic prompt-based watermarking framework for LLM-generated text that operates without access to model weights or decoding internals. The pipeline was evaluated across 25 open-source LLMs including fine-tuned and distilled variants trained using LoRA and DeepSpeed. Achieves 88%+ detection accuracy under paraphrasing, fine-tuning, and adversarial prompt attacks.

**SSRIW (2023)** — <https://github.com/cent664/SSRIW>

A ViT-based cross-attention watermarking system that directly trains an invariant domain alongside the watermarking task using a triplet loss objective. The cross-attention mechanism localizes semantically relevant regions for watermark embedding, and the invariant domain provides robustness under compound distortions. Achieves state-of-the-art comparable accuracy under common noise types with improved performance under combined augmentations.

**PTL (2022)** — [https://github.com/kcnishan/Perspective\\_Transformation\\_Layer](https://github.com/kcnishan/Perspective_Transformation_Layer)

A lightweight differentiable layer that learns 2D projections of 3D views for multiple viewpoints, implemented as a drop-in module within standard deep learning pipelines. Rather than relying on fixed augmentation pipelines, PTL internalizes perspective warping as a learnable operator, allowing the model to adaptively select viewpoint transformations that improve downstream task accuracy.

## TECHNICAL SKILLS

---

|                     |                                                             |
|---------------------|-------------------------------------------------------------|
| <b>ML Concepts:</b> | Watermarking, Computer Vision, NLP, LLMs                    |
| <b>Frameworks:</b>  | PyTorch, TensorFlow, HuggingFace Transformers, scikit-learn |
| <b>Programming:</b> | Python, Bash, CUDA, Git, Linux                              |
| <b>Distributed:</b> | DeepSpeed, LoRA, PyTorch DDP, Slurm/HPC                     |
| <b>Agentic AI:</b>  | LangGraph, RAG, LLM-as-Judge, ChromaDB                      |
| <b>Deployment:</b>  | FastAPI, OpenAI & Anthropic APIs                            |

## REFERENCES

---

**Dr. Xin Zhong**  
Associate Professor  
University of Nebraska Omaha  
✉ [xzhong@unomaha.edu](mailto:xzhong@unomaha.edu)  
☎ +1 (402) 554-6637  
PKI 281A, Omaha, NE 68182

**Dr. Qiuming Zhu**  
Professor  
University of Nebraska Omaha  
✉ [qzhu@unomaha.edu](mailto:qzhu@unomaha.edu)  
☎ +1 (402) 554-3685  
PKI 281D, Omaha, NE 68182

**Dr. Dario Gheri**  
Associate Professor  
University of Nebraska Omaha  
✉ [dghersi@unomaha.edu](mailto:dghersi@unomaha.edu)  
☎ +1 (402) 554-6004  
PKI 173B, Omaha, NE 68182